

SQL Programming and Normalization

Colton Blackwell
CMPT 354
301451278

DBMS environment used: mySQL Workbench

Connection Details
Name: local host
Host: 127.0.0.1
Port: 3306
Login User: root
Current User: root@localhost
SSL cipher: TLS_AES_256_GCM_SHA384
Server
Product: MySQL Community Server - GPL
Version: 8.0.35
Connector
Version: C++ 8.1.0

Part 1: SQL Programming

Task 1: Create Table statements (employee, types, sales)

employee

Table: **employee**

Columns:

eid	int PK
name	varchar(20)
salary	int
dept	varchar(20)

Table:	employee
Create Table:	<pre>CREATE TABLE `employee` (`eid` int NOT NULL, `name` varchar(20) DEFAULT NULL, `salary` int DEFAULT NULL, `dept` varchar(20) DEFAULT NULL, PRIMARY KEY (`eid`), CONSTRAINT `employee_chk_1` CHECK (((`salary` >= 5000) and (`salary` <= 20000)))) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci</pre>

*Used VARCHAR(20) instead of CHAR(20)

types

Table: **types**

Columns:

item varchar(20) PK
color varchar(20) PK

Table:

types

Create Table:

```
CREATE TABLE `types` (  
  `item` varchar(20) NOT NULL,  
  `color` varchar(20) NOT NULL,  
  PRIMARY KEY (`item`,`color`)  
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci
```

sales

Table: **sales**

Columns:

dept varchar(20) PK
item varchar(20) PK

Table:

sales

Create Table:

```
CREATE TABLE `sales` (  
  `dept` varchar(20) NOT NULL,  
  `item` varchar(20) NOT NULL,  
  PRIMARY KEY (`dept`,`item`)  
) ENGINE=InnoDB DEFAULT CHARSET=utf8mb4 COLLATE=utf8mb4_0900_ai_ci
```

Task 2: Insert records into the tables

employee

```
17:01:06 INSERT INTO employee (eid, name, salary, dept) VALUES (666,"Hoffman", 16000, "Cosmetics")
17:01:06 INSERT INTO employee (eid, name, salary, dept) VALUES (555,"Nelson", 6000, "Toy")
17:01:06 INSERT INTO employee (eid, name, salary, dept) VALUES (444,"Lewis", 12000, "Stationary")
17:01:06 INSERT INTO employee (eid, name, salary, dept) VALUES (333,"Morgan", 10000, "Cosmetics")
17:01:06 INSERT INTO employee (eid, name, salary, dept) VALUES (222,"Anderson", 8000, "Toy")
17:01:06 INSERT INTO employee (eid, name, salary, dept) VALUES (111,"Jane", 8000, "Household")
```

	eid	name	salary	dept
▶	111	Jane	8000	Household
	222	Anderson	8000	Toy
	333	Morgan	10000	Cosmetics
	444	Lewis	12000	Stationary
	555	Nelson	6000	Toy
	666	Hoffman	16000	Cosmetics

* Row with name = 'Peter' was added in Question 7

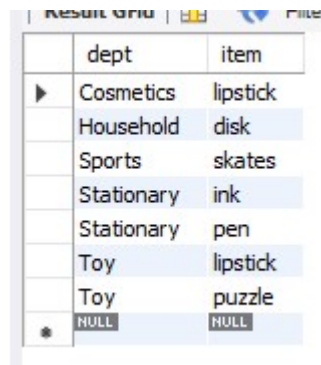
types

```
17:08:19 INSERT INTO types VALUES ("ink","blue")
17:08:19 INSERT INTO types VALUES ("ink","red")
17:08:19 INSERT INTO types VALUES ("puzzle","black")
17:08:19 INSERT INTO types VALUES ("pen","black")
17:08:19 INSERT INTO types VALUES ("lipstick","red")
17:08:19 INSERT INTO types VALUES ("pen","red")
```

	item	color
▶	ink	blue
	ink	red
	lipstick	red
	pen	black
	pen	red
	puzzle	black
*	NULL	NULL

Sales

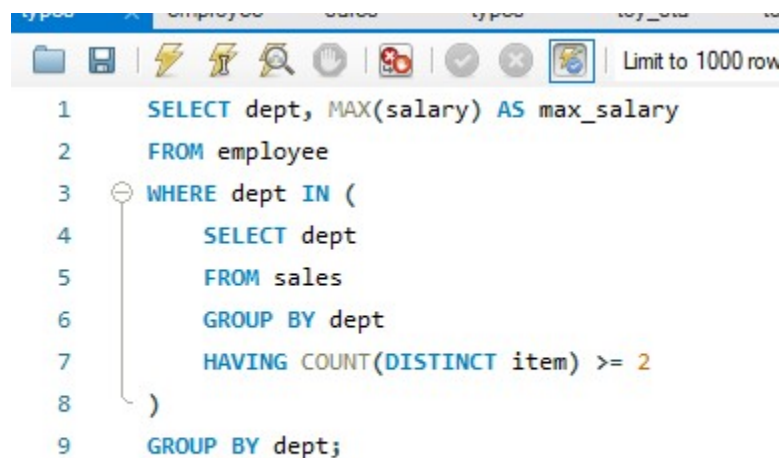
```
17:05:09 INSERT INTO sales VALUES ("Toy", "lipstick")
17:05:09 INSERT INTO sales VALUES ("Sports", "skates")
17:05:09 INSERT INTO sales VALUES ("Household", "disk")
17:05:09 INSERT INTO sales VALUES ("Stationary", "ink")
17:05:09 INSERT INTO sales VALUES ("Toy", "puzzle")
17:05:09 INSERT INTO sales VALUES ("Cosmetics", "lipstick")
17:05:09 INSERT INTO sales VALUES ("Stationary", "pen")
```



	dept	item
▶	Cosmetics	lipstick
	Household	disk
	Sports	skates
	Stationary	ink
	Stationary	pen
	Toy	lipstick
	Toy	puzzle
*	NULL	NULL

Task 3: Query Data for Specific instances

1. Compute the maximum salary for each department that sells at least two distinct items.



```
1  SELECT dept, MAX(salary) AS max_salary
2  FROM employee
3  WHERE dept IN (
4      SELECT dept
5      FROM sales
6      GROUP BY dept
7      HAVING COUNT(DISTINCT item) >= 2
8  )
9  GROUP BY dept;
```

Result Grid		
	dept	max_salary
▶	Toy	8000
	Stationary	12000

2. Compute the names of the employees who work in a department that sells some item in black colour.

```

1 • SELECT DISTINCT e.name
2   FROM employee e
3   WHERE e.dept IN (
4       SELECT DISTINCT s.dept
5       FROM sales s
6       WHERE s.item IN (
7           SELECT DISTINCT t.item
8           FROM types t
9           WHERE t.color = 'black'
10      )
11  );

```

Result Grid	
	name
▶	Anderson
	Lewis
	Nelson

3. For each department that has a larger average salary than that of "Stationary" department, find its average salary.

```
types x
Limit to 1000 rows
1 • SELECT e.dept, AVG(e.salary) AS avg_salary
2 FROM employee e
3 WHERE e.dept <> "Stationary"
4 GROUP BY e.dept
5 HAVING AVG(e.salary) > (SELECT AVG(salary) FROM employee WHERE dept = "Stationary");
```

Result Grid | Filter Rows:

	dept	avg_salary
▶	Cosmetics	13000.0000

4. Find the number of the departments that have a smaller average salary than that of "Stationary" department.

```
types x employee sales types
Limit to 1000 rows
1 • SELECT e.dept, COUNT(e.dept) AS num_departments
2 FROM employee e
3 GROUP BY e.dept
4 HAVING AVG(e.salary) < (SELECT AVG(salary) FROM employee WHERE dept = 'Stationary');
```

Result Grid | Filter Rows:

	dept	num_departments
▶	Household	1
	Toy	2

5. Which department pays every of its employees at least 7000?

```
types x employee sales types
Limit to 1000 rows
1 • SELECT dept
2 FROM employee
3 GROUP BY dept
4 HAVING MIN(salary) >= 7000;
```

Result Grid		Filter R
	dept	
▶	Household	
	Cosmetics	
	Stationary	

6. Which department sells all items sold by 'Cosmetics' department?

```

1  SELECT DISTINCT s1.dept
2  FROM sales s1
3  WHERE NOT EXISTS (
4      SELECT item
5      FROM sales s2
6      WHERE s2.dept = 'Cosmetics'
7      AND NOT EXISTS (
8          SELECT *
9          FROM sales s3
10         WHERE s3.dept = s1.dept AND s3.item = s2.item
11     )
12 );

```

	dept
▶	Cosmetics
	Toy

Part 2: FD and Normalization

Question 1: Anomalies in Relation R with FD's F

Relation below is used for example purposes:

S	A	C	D	T
Johnny	324 Redwood St	Math	Math	Prof. Hughes
Johnny	324 Redwood St	History	History	Prof. Petterson
Max	929 Cedar St	Science	Science	Prof. Kuzmenko
Ted	223 Oak St	Math	Math	Prof. Hughes

Data Redundancy:

The address is repeated for each course a student takes, leading to redundancy

- Ex. Student Johnny has the same address for Math and History courses.

Update anomaly:

If a student updates their address, we need to update multiple rows to maintain consistency.

- Ex. If Student Johnny moves from 324 Redwood St to a different address, we need to update that for both occurrences of his Math and History courses.

Insertion Anomaly:

To insert a new course, we have to provide information about a teacher and department due to functional dependencies. If a department does not have any teachers at the moment, we cant insert info about a new course for that department until a teacher is assigned.

Deletion Anomaly:

If a student drops all their courses, we lose info about their address. Additionally, if a teacher stops teaching all courses, we lose information about their department. Deleting a department also deletes information about all the teachers in that department.

Question 2:

Q2 (1): Find all keys of R with respect to F.

If a key must uniquely determine all oher attributes in the relation R, and given...

$R = (S, A, C, D, T)$

With functional dependencies...

C -> T
S -> A
T -> D

{S, C} does not occur in RHS of any FD which means that they must be a part of every key

Key Computation:
 $SC^+ = \{SACDT\}$

\therefore SC is the only key!

Q2 (2): Test if R in BCNF with respect to F, why?

Example of redundancy:

Student	Address	Course	Department	Teacher
Johnny	324 Redwood St	MACM 101	Math	Prof. Hughes
Johnny	324 Redwood St	MACM 101	Math	Prof. Hughes
Max	929 Cedar St	MACM 101	Math	Prof. Hughes
Max	929 Cedar St	MACM 101	Math	Prof. Hughes

Student -> Address

Course -> Teacher

Teacher -> Department

Based on results from Question 2 (1), we know that the only key is {Student, Course}.
In each FD, the left side is not a superkey.

Ex. In Student -> Address, Student is not a superkey therefore there are multiple instances of redundant data for Johnny -> 324 Redwood St

Q2 (3): Produce a BCNF decomposition through a series of binary decomposition.

We start with reviewing the data that we are given

R(Student, Address, Course, Department, Teacher)

Course -> Teacher

Student -> Address

Teacher -> Department

Key = {Student, Course}

We then pick any BCNF violation that we observe in F. In this case, they are all violations on BCNF. We choose Course because it is not the candidate key.

$\{Course\}^+ = \{Course, Teacher, Department\}$

Decompose R into...

1. $R_1 = X + (\text{all attributes determined by } X)$
2. $R_2 = (R - X) + X$ (all attributes not determined by X, plus X)

Relation1(Course, Teacher, Department)

Relation2(Course, Address, Student)

For relation1 FD's are implied that course->Teacher and Teacher-> department. Course is the superkey in relation1 so relation1 is in BCNF.

For relation2 the only FD is Student -> Address so we repeat decomposition on R

$Student^+ = \{Student, Address\}$

(new)Relation2(Student, Address)

Relation3(Student, Course)

Relation3 holds for BCNF as Student is superkey. The final decomposition is composed of sub-relations that all their functional dependencies hold BCNF.

New and revised Relations that hold for BCNF

Relation1		
Course	Teacher	Department
MACM 101	Prof. Hughes	Math
CMPT 125	Prof. Petterson	Computer Science
CMPT 321	Prof. Tanev	Computer Science

Relation2	
Student	Address
Elias	123 Cedar Road
Quinn	221 Pine ave.
Alex	252 Thorne blvd

Relation3	
Student	Course
Elias	MACM 101
Quinn	CMPT 321
Alex	MACM 101

Q2 (4): Why are the new decomposed tables in a better representation?

The new decomposed tables are in a better representation as they reduce anomalies as shown in Part 2 Question 1. These relations are more specialized in fulfilling their designated roles, effectively minimizing data redundancy. The new relations address deletion anomalies, ensuring that removing a student's record does not accidentally delete other attributes, such as courses from the table.

Relation1 provides insights into teachers, their associated departments, and the course they instruct.

Relation2 details information about students and their respective addresses.

Relation3 specifically outlines the students and the courses which they are currently enrolled in.

Q2 (5): Is the final decomposition in 3 dependency-preserving ?

Yes the final decomposition is dependency preserving as all the FD's remained in their respective separate tables.

Relation2: Student -> Address

Relation1: Course -> Teacher

Teacher -> Department

Q2 (6): Is the original schema R in 3NF with respect to F?

No the original schema is not in 3NF because there is a transitive dependency. Additionally, SC is the only key which means C, S, T are not superkeys and T, A, D are not prime

$C \rightarrow T$ and $T \rightarrow D$

The non-prime D depends on non-prime T

Question 2 (7): Produce a 3NF decomposition that is lossless and dependency-preserving.

Consider $R = \{SACDT\}$ and $F = \{C \rightarrow T, S \rightarrow A, T \rightarrow D\}$

All FD's are already in minimal cover and are not redundant.

For each FD $X \rightarrow A$ in the minimal cover, add a relation $R_i = XA$.

$R_1 = CT$

$R_2 = SA$

$R_3 = TD$

No key is contained in any R_i , so the addition of the key in decomposition is mandatory.

$R_4 = SC$

3NF Decomposition:

$\{CT, SA, TD, SC\}$

Question 2 (8): Is the decomposition produced in 7 in BCNF?

The outcome is in BCNF as the left-hand side of each relation has a superkey within their respective relation.